


Changing evidence: predicting failure does not identify the remedy

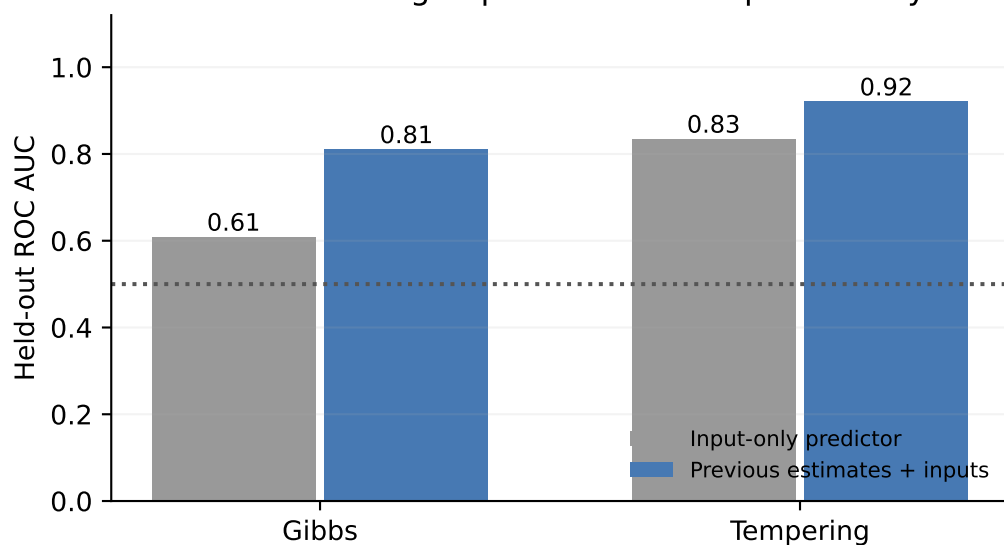
Reversing evidence: Gibbs, 4 sweeps/query



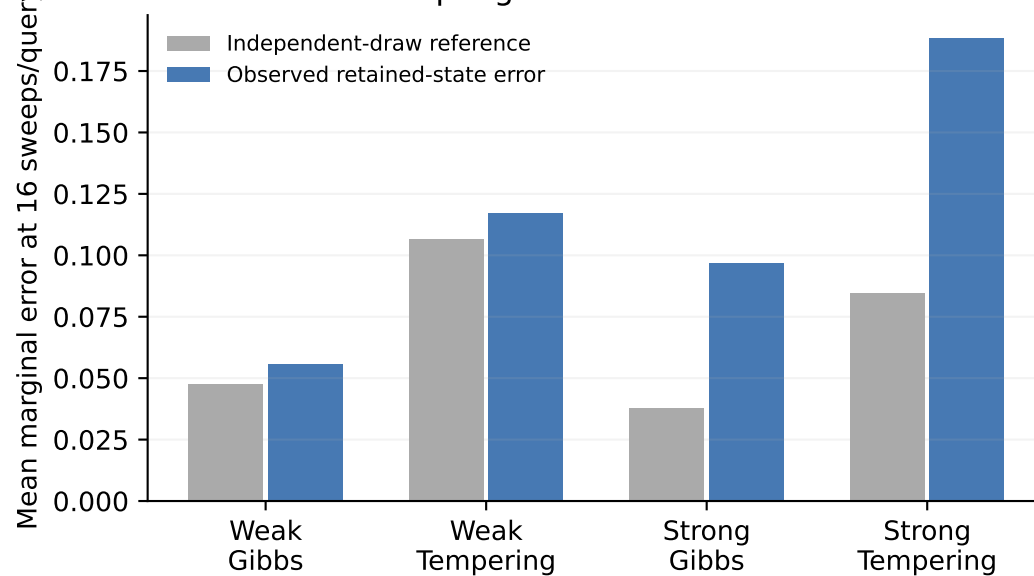
Retained state adds history dependence



Failure ranking improves with sampler history



Post-hoc: sampling noise versus additional error



CPU software simulation on 12-spin models; 8 held-out seeds. Top row: strong coherent round trips.
Predictive AUC pools paired conditions descriptively. The independent-draw reference is not a universal lower bound.